

TASK 1 - FDA-Grade Data Provenance ETL Pipeline

1.1 Summary

- Design an end-to-end ETL pipeline to:
- Ingest clinical text from multiple sources
- Validate schema at every stage
- Scrub PHI using formal rules
- Track full provenance (timestamps, rules, transformations)
- Output QLM-ready datasets
- Expose /provenance API for lineage retrieval

1.2 Inputs

- Raw clinical text files (CSV, JSON, HL7)
- Metadata describing each source (file name, schema version, hash)
- PHI redaction rules (names, dates, phone numbers, addresses)
- Input and output schema definitions (JSON or Parquet)

1.3 Expected Outputs

- QLM-ready dataset (cleaned, PHI redacted, hashed)
- Provenance logs (per record / per transformation)
- Integrity verification (SHA-256 hash per batch or record)
- /provenance REST API
- Developer and compliance documentation

1.4 Technical Implementation

| Component | Azure Service / Tool | Notes |
|------------------------|---|--|
| Raw & Curated Storage | Azure Data Lake Storage Gen2 | Hierarchical namespace, encrypted, immutable |
| ETL Processing | Azure Databricks (PySpark notebooks) | Multi-source ingestion & transformation |
| Multi-source Ingestion | PySpark CSV/JSON reader + HL7 parser (hl7apy) | Handles large datasets |
| Schema Validation | jsonschema in Databricks | Validates input/output schema |
| PHI Redaction | Python regex / NLP rules | Names → <NAME_REDACTED>; Phones → <PHONE_REDACTED> |
| Provenance Logging | JSON logs in /audit/provenance/ (ADLS) | Append-only, immutable |
| Integrity Verification | SHA-256 hashes in /integrity/hashes.json | Ensures data integrity |
| REST API | FastAPI on Azure App Service / Functions | Exposes /provenance/{record_id} |
| RBAC | Azure RBAC | Enforces least-privilege access |
| Monitoring & Audit | Azure Monitor + WORM storage | Tracks access & modifications for compliance |

1.5 Example ETL Code (Databricks / PySpark)

```
from pyspark.sql import SparkSession  
  
import hashlib, re, json  
  
from datetime import datetime  
  
spark = SparkSession.builder.appName("FDA_ETL").getOrCreate()
```

```
def ingest_file(path, source_type):  
  
    if source_type == "csv":  
  
        df = spark.read.csv(path, header=True)  
  
    elif source_type == "json":
```

```
df = spark.read.json(path)

elif source_type == "hl7":
    df = parse_hl7(path) # Custom HL7 parser

return df


def redact_phi(text):
    text = re.sub(r"\b\d{3}[-.]\d{3}[-.]\d{4}\b", "<PHONE_REDACTED>", text)
    text = re.sub(r"\b([A-Z][a-z]+ [A-Z][a-z]+)\b", "<NAME_REDACTED>", text)

    return text


def compute_hash(text):
    return hashlib.sha256(text.encode("utf-8")).hexdigest()

def log_provenance(record_id, step, input_hash, output_hash):
    entry = {
        "record_id": record_id,
        "timestamp": datetime.utcnow().isoformat(),
        "transformation": step,
        "input_hash": input_hash,
        "output_hash": output_hash
    }

    with open(f"/dbfs/mnt/datalake/audit/provenance/{record_id}.json", "w") as f:
        f.write(json.dumps(entry))
```

1.6 Provenance REST API

```
from fastapi import FastAPI
import json, os
app = FastAPI()
@app.get("/provenance/{record_id}")
def get_provenance(record_id: str):
    file_path = f"/mnt/datalake/audit/provenance/{record_id}.json"
    if os.path.exists(file_path):
        with open(file_path) as f:
            return json.load(f)
    return {"error": "Record not found"}
```

1.6.1 Sample /provenance API Output

```
{
    "record_id": "rec_001",
    "timestamp": "2025-02-01T12:44:21Z",
    "transformation": "PHI_redaction",
    "input_hash": "a89fd1d3b97d0ba9ac...",
    "output_hash": "19bd9f1170ffbac120..."
}
```

1.7 Lineage Architecture Diagram

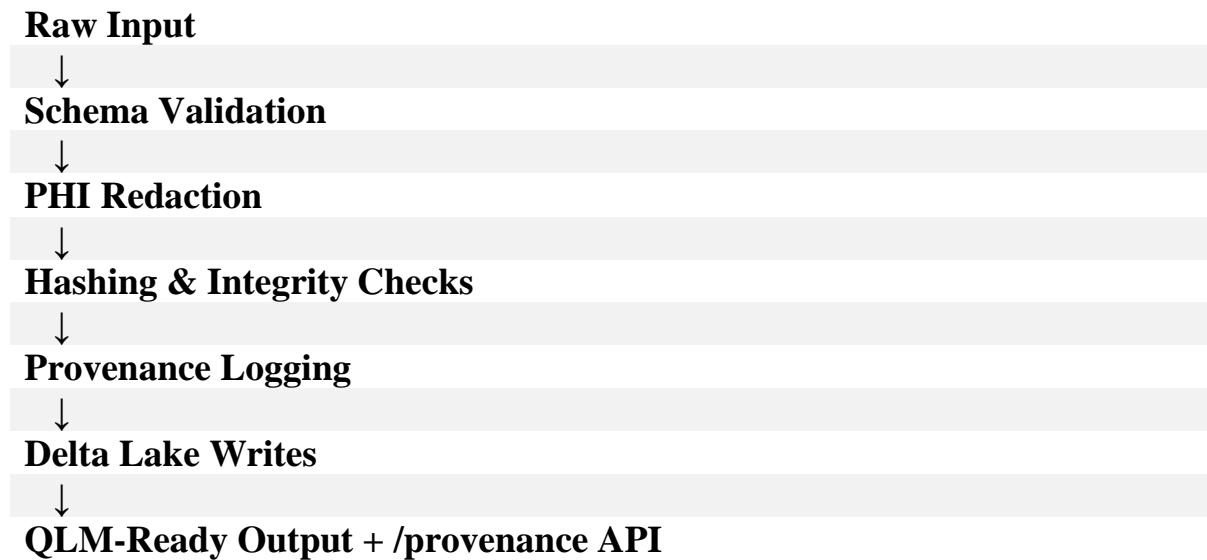
[Source 1: Blob JSON] --> [ADF Copy] --> [Synapse Data Flow:
Validate/Scrub/Hash] --> [Parquet Sink: ADLS QLM-Ready]

[Source 2: CSV] |

v

[SQL: Provenance Logs + Hashes] <--> [Purview: Auto-Lineage Scan] <-->
[Function API: /provenance]

1.8 Logical Architecture Diagram



1.9 Deployment & Configuration Guide

Databricks

1. Create service principal and mount ADLS Gen2 with OAuth2.
2. Upload ETL notebooks for ingestion, validation, PHI redaction, hashing.
3. Create Databricks Jobs for scheduled batch processing.
4. Enable cluster-level secure networking (no public IP).

App Service

1. Containerize the FastAPI app using Docker.
2. Deploy to Azure App Service or Azure Function App.
3. Use Managed Identity to access ADLS /audit/provenance files.
4. Apply Azure Front Door or APIM for secure access control.

Purview + Monitoring

1. Register ADLS, Databricks, App Service in Purview.
2. Enable diagnostic logs (Storage Analytics + Log Analytics).
3. Configure alerts for unusual access patterns.

TASK 2 - Tiered, Regulated Data Lake Architecture

2.1 Summary

Design a secure, multi-zone data lake:

Raw Zone – exact copy of data, immutable, encrypted

Curated Zone – standardized, partially PHI redacted, versioned

QLM-Ready Zone – fully PHI redacted, hashed, minimal normalized fields

2.2 Storage Zones & Azure Mapping

| Zone | Description | Azure Implementation |
|-------------|--|-----------------------------|
| Raw | Exact copy, immutable, encrypted | ADLS Gen2 + WORM |
| Curated | Standardized, validated, partial PHI removal | ADLS Gen2 + Delta Lake |
| QLM-Ready | Fully PHI redacted, hashed | ADLS Gen2 + Delta Lake |

2.3 RBAC

| Role | Raw Zone | Curated Zone | QLM-Ready Zone | Logs | API |
|--------------------|-----------------|---------------------|-----------------------|-------------|------------|
| Data Engineer | ✓ | ✓ | ✓ | ✓ | ✓ |
| Data Scientist | ✗ | ✓ | ✓ | ✗ | ✗ |
| Compliance Officer | ✓ | ✓ | ✗ | ✓ | ✓ |
| Analyst | ✗ | ✗ | ✓ | ✗ | ✗ |

2.4 Versioning & Audit

Delta Lake tables on ADLS Gen2 (Curated & QLM zones)

Every write produces: commit logs, snapshots, time travel

Audit logs track: access, modifications, PHI redaction version, pipeline version

Logs stored in immutable WORM storage

2.4.1 Example Versioned Dataset

Describe History curated.clinical_notes;

| Version | Timestamp | Operation |
|----------------|---------------------|------------------------|
| 0 | 2025-02-01 10:00:00 | CREATE TABLE |
| 1 | 2025-02-01 11:30:12 | MERGE (batch_20250201) |
| 2 | 2025-02-02 09:10:44 | WRITE (batch_20250202) |

2.4.2 Time Travel Query Example

```
SELECT * FROM curated.clinical_notes VERSION AS OF 1;
```

2.4.3 Example Audit Logs

Access Log

```
{  
  "user": "data_engineer",  
  "action": "READ",  
  "dataset": "curated.clinical_notes",  
  "timestamp": "2025-02-03T16:11:20Z",  
  "ip": "10.0.2.14"  
}
```

Transformation Log

```
{  
  "pipeline": "ETL_v3",  
  "transformation": "PHI_redaction_v1",  
  "timestamp": "2025-02-03T09:14:11Z",  
  "records_processed": 14592  
}
```

2.5 Compliance Mapping

| Regulation | How Solution Meets It |
|-------------|---|
| HIPAA | PHI removal, RBAC, encryption |
| FDA | Full data lineage, integrity checks, reproducible pipelines |
| CFR Part 11 | Timestamped transformations, immutable logs |
| Security | Encryption, RBAC, network segmentation |

2.5.1 Expanded Compliance Mapping

HIPAA Compliance

- Full PHI removal in the QLM zone.
- RBAC ensures only authorized users access raw/curated zones.
- Encryption in transit (TLS1.2+) and at rest (Azure SSE).
- Access logs stored in immutable WORM storage.

FDA Requirements

- Complete record-level lineage via provenance logs.
- SHA-256 integrity hashing for tamper detection.
- Reproducible ETL via versioned notebooks and config files.
- Immutable WORM audit logs meet FDA documentation rules.

CFR Part 11

- Immutable audit trail via WORM + Delta commit logs.
- Timestamped entries for all transformations.
- Identity-bound changes tracked via Azure AD.
- Electronic signatures through RBAC-controlled access.
- No silent modification of data (Delta Lake ACID guarantees).

2.6 Deliverables

Architecture Diagrams

Physical (Text; PNG in repo):

[Azure Portal] --> [ADLS Gen2 Account]

|

+--> [Containers: raw | curated | qlm-ready (Delta Tables)]

|

[Synapse Workspace] <--> [ADF Pipelines] <--> [Purview Catalog]

|

v

[Log Analytics] <-- [RBAC Roles] <-- [Monitor Alerts]